

ML LAB WEEK 14

CNN

Name-Nisschay Khandelwal

SRN-PES2UG23CS394

1. Introduction:

This lab focused on building, training, and testing a Convolutional Neural Network (CNN) to classify images of hands playing Rock, Paper, or Scissors. The objective was to implement a complete deep learning pipeline including data preprocessing, model architecture design, training loop implementation, and evaluation on unseen test data. The dataset was obtained from Kaggle and split into 80% training and 20% testing sets.

2. Model Architecture:

The CNN architecture consists of two main components: a convolutional block and a fully-connected classifier. The convolutional block contains three sequential layers, each consisting of a Conv2d layer, ReLU activation, and MaxPool2d layer. The first convolutional layer takes 3 input channels (RGB) and outputs 16 channels using a 3x3 kernel with padding of 1. The second layer increases channels from 16 to 32, and the third layer from 32 to 64, both using the same kernel size and padding. Each convolutional layer is followed by a 2x2 max pooling operation, which reduces the spatial dimensions by half at each stage. Starting from 128x128 input images, the feature maps are progressively reduced to 64x64, then 32x32, and finally 16x16.

The fully-connected classifier begins with a Flatten layer that converts the 64x16x16 feature maps into a 1D vector of 16,384 elements. This is followed by a Linear layer that reduces dimensionality to 256 neurons, a ReLU activation function, and a Dropout layer with probability 0.3 to prevent overfitting. Finally, a Linear layer maps the 256 features to 3 output classes corresponding to rock, paper, and scissors.

3. Training and Performance:

The model was trained using the Adam optimizer with a learning rate of 0.001. CrossEntropyLoss was used as the loss function, which is standard for multi-class classification problems. The training was conducted for 10 epochs with a batch size of 32 for both training and testing data loaders. The training data loader was shuffled to ensure random sampling during training, while the test loader was not shuffled to maintain consistency in evaluation. The final test accuracy achieved by the model would be displayed after running the evaluation cell in the notebook.

4. Conclusion and Analysis:

The CNN model successfully learned to classify hand gestures into rock, paper, and scissors categories. The architecture proved effective for this image classification task, with the convolutional layers extracting relevant features and the fully-connected layers performing accurate classification. One challenge faced during implementation was ensuring proper data preprocessing with correct normalization parameters and image resizing to maintain aspect ratios.

To improve the model's accuracy in the future, data augmentation techniques such as random rotations, flips, and brightness adjustments could be applied to increase the diversity of training samples and improve generalization. Additionally, experimenting with deeper architectures or using transfer learning with pre-trained models like ResNet or EfficientNet could potentially boost performance further. Fine-tuning hyperparameters such as learning rate scheduling or increasing the number of epochs with early stopping could also lead to better results.

SCREENSHOTS

```
# Install required packages
!pip install kagglehub --quiet
!pip install Pillow numpy --quiet

import kagglehub
path = kagglehub.dataset_download("drgfreeman/rockpaperscissors")
print("Path to dataset files:", path)

Using Colab cache for faster access to the 'rockpaperscissors' dataset.
Path to dataset files: /kaggle/input/rockpaperscissors

import shutil
import os

src_root = "/kaggle/input/rockpaperscissors"
dst_root = "/content/dataset"

os.makedirs(dst_root, exist_ok=True)
```

```
Using Colab cache for faster access to the 'rockpaperscissors' dataset.
Path to dataset files: /kaggle/input/rockpaperscissors

import shutil
import os

src_root = "/kaggle/input/rockpaperscissors"
dst_root = "/content/dataset"

os.makedirs(dst_root, exist_ok=True)

folders_to_copy = ["rock", "paper", "scissors"]

for folder in folders_to_copy:
    src_path = os.path.join(src_root, folder)
    dst_path = os.path.join(dst_root, folder)

    if os.path.exists(src_path):
        shutil.copytree(src_path, dst_path, dirs_exist_ok=True)
        print("Copied: " + folder)
    else:
        print("Folder not found: " + folder)

Copied: rock
Copied: paper
Copied: scissors

Step 2: Imports and Device Setup

Import the necessary libraries and check if a GPU is available.
```

```
import os
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
```

boilerplate.ipynb - Colab

https://colab.research.google.com/drive/1CtjhXzFLYLRNwAxtSllrvpjPeeg5fLV#scrollTo=13-RXEbuuxu

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Step 2: Imports and Device Setup

Import the necessary libraries and check if a GPU is available.

```
[14] On
import os
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader, random_split
from PIL import Image
import numpy as np

# TODO: Set the 'device' variable
# Check if CUDA (GPU) is available, otherwise use CPU
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

print("Using device:", device)
... Using device: cuda
```

Step 3: Data Loading and Preprocessing

Here we will define our image transformations, load the dataset, split it, and create DataLoaders.

```
[15] On
DATA_DIR = "/content/dataset"

# TODO: Define the image transforms
# We need to:
# 1. Resize all images to 128x128
# 2. Convert them to Tensors
# 3. Normalize them (mean=0.5, std=0.5)
transform = transforms.Compose([
    transforms.Resize((128, 128)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
])
```

Variables Terminal

boilerplate.ipynb - Colab

https://colab.research.google.com/drive/1CtjhXzFLYLRNwAxtSllrvpjPeeg5fLV#scrollTo=13-RXEbuuxu

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

```
[16] On
# We need to:
# 1. Resize all images to 128x128
# 2. Convert them to Tensors
# 3. Normalize them (mean=0.5, std=0.5)
transform = transforms.Compose([
    transforms.Resize((128, 128)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
])

# Load dataset using ImageFolder
full_dataset = datasets.ImageFolder(DATA_DIR, transform=transform)

class_names = full_dataset.classes
print("Classes:", class_names)

# TODO: Split the dataset
# We want 80% for training and 20% for testing
train_size = int(0.8 * len(full_dataset))
test_size = len(full_dataset) - train_size

# TODO: Use random_split to create train_dataset and test_dataset
train_dataset, test_dataset = random_split(full_dataset, [train_size, test_size])

# TODO: Create the DataLoaders
# Use a batch_size of 32
# Shuffle the training loader, but not the test loader
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

print(f"Total Images: {len(full_dataset)}")
print(f"Training Images: {len(train_dataset)}")
print(f"Test Images: {len(test_dataset)}")
```

... Classes: ['paper', 'rock', 'scissors']
Total Images: 7100
Training Images: 5680
Test Images: 1420

Variables Terminal

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

Step 4: Define the CNN Model

Fill in the `conv_block` and `fc_block` with the correct layers.

```
class RPS_CNN(nn.Module):
    def __init__(self):
        super(RPS_CNN, self).__init__()

        # TODO: Define the convolutional block
        # We want 3 blocks:
        # 1. Conv2d(16 -> 32 channels, kernel=3, padding=1), ReLU, MaxPool2d(2)
        # 2. Conv2d(32 -> 64 channels, kernel=3, padding=1), ReLU, MaxPool2d(2)
        # 3. Conv2d(64 -> 128 channels, kernel=3, padding=1), ReLU, MaxPool2d(2)
        self.conv_block = nn.Sequential(
            nn.Conv2d(1, 16, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2),
            nn.Conv2d(16, 32, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2),
            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2)
        )

        # After 3 MaxPool(2) layers, our 128x128 image becomes:
        # 128 -> 64 -> 32 -> 16
        # So the flattened size is 64 * 16 * 16      # TODO: Define the fully-connected (classifier) block
        # We want:
        # 1. Flatten the input
        # 2. Linear layer (64 * 16 * 16 -> 256)
        # 3. ReLU
        # 4. Dropout (p=0.3)
        # 5. Linear layer (256 -> 3) (3 classes: rock, paper, scissors)
        self.fc = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64 * 16 * 16, 256),
            nn.ReLU(),
            nn.Dropout(p=0.3),
            nn.Linear(256, 3),
            nn.ReLU()
        )

    def forward(self, x):
        x = self.conv_block(x)
        x = self.fc(x)
        return x
```

Variables Terminal ✓ 5:55PM T4 (Python 3)

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

```
# TODO: Initialize the model, criterion, and optimizer
# 1. Create an instance of RPS_CNN and move it to the 'device'
model = RPS_CNN().to(device)

# 2. Define the loss function (Criterion). Use CrossEntropyLoss for classification.
criterion = nn.CrossEntropyLoss()

# 3. Define the optimizer. Use Adam with a learning rate of 0.001
optimizer = optim.Adam(model.parameters(), lr=0.001)

print(model)
```

```
... RPS_CNN(
    (conv_block): Sequential(
        (0): Conv2d(1, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
        (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (3): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (4): ReLU()
        (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (6): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (7): ReLU()
        (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (fc): Sequential(
        (0): Flatten(start_dim=1, end_dim=-1)
        (1): Linear(in_features=16384, out_features=256, bias=True)
        (2): ReLU()
    )
)
```

Variables Terminal ✓ 5:55PM T4 (Python 3)

boilerplate.ipynb - Colab

https://colab.research.google.com/drive/1CtjhXzFLYRNwAxISlrlvpjPeeg5fLV#scrollTo=wexPK8V3y3fx

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Step 5: Train the Model

Fill in the core training steps inside the loop.

```
EPOCHS = 10

for epoch in range(EPOCHS):
    model.train() # Set the model to training mode
    total_loss = 0

    for images, labels in train_loader:
        # Move data to the correct device
        images, labels = images.to(device), labels.to(device)

        # TODO: Implement the training steps
        # 1. Clear the gradients (optimizer.zero_grad())
        optimizer.zero_grad()

        # 2. Perform a forward pass (get model outputs)
        outputs = model(images)

        # 3. Calculate the loss (using criterion)
        loss = criterion(outputs, labels)

        # 4. Perform a backward pass (loss.backward())
        loss.backward()

        # 5. Update the weights (optimizer.step())
        optimizer.step()

        total_loss += loss.item()

    print(f"Epoch {epoch+1}/{EPOCHS}, loss = {total_loss/len(train_loader):.4f}")
print("Training complete!")

--- Epoch 1/10, loss = 0.0066
Epoch 2/10, loss = 0.0074
```

Variables Terminal ✓ 5:55PM T4 (Python 3)

boilerplate.ipynb - Colab

https://colab.research.google.com/drive/1CtjhXzFLYRNwAxISlrlvpjPeeg5fLV#scrollTo=wexPK8V3y3fx

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Step 6: Evaluate the Model

Test the model's accuracy on the unseen test set.

```
model.eval() # Set the model to evaluation mode
correct = 0
total = 0

# TODO: Use torch.no_grad()
# we don't need to calculate gradients during evaluation
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)

        # TODO: Get model predictions
        # 1. Get the raw model outputs (logits)
        outputs = model(images)

        # 2. Get the predicted class (the one with the highest score)
        # Hint: use torch.max(outputs, 1)
        _, predicted = torch.max(outputs, 1)

        total += labels.size(0)
```

Variables Terminal ✓ 5:55PM T4 (Python 3)

The screenshot shows a Google Colab notebook titled "boilerplate.ipynb". The code cell (cell 18) contains the following Python code:

```
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)

        # TODO: Get model predictions
        # 1. Get the raw model outputs (logits)
        outputs = model(images)

        # 2. Get the predicted class (the one with the highest score)
        # Hint: use torch.max(outputs, 1)
        _, predicted = torch.max(outputs, 1)

        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f"Test Accuracy: {(100 * correct / total):.2f}%")
Test Accuracy: 95.89%
```

The notebook also includes a section titled "Step 7: Test on a Single Image" with the following code:

```
def predict_image(model, img_path):
    model.eval()

    img = Image.open(img_path).convert("RGB")
    # Apply the same transforms as training, and add a batch dimension (unsqueeze)
    img = transform(img).unsqueeze(0).to(device)

    with torch.no_grad():
        # TODO: Get the model prediction
        # 1. Get the raw model outputs (logits)
        output = model(img)

        # 2. Get the predicted class index
        _, pred = torch.max(output, 1)
```

The screenshot shows the same Google Colab notebook. The code cell (cell 18) has been run, and the output is displayed:

```
correct += (predicted == labels).sum().item()
print(f"Test Accuracy: {(100 * correct / total):.2f}%")
Test Accuracy: 95.89%
```

The notebook also includes a section titled "Step 7: Test on a Single Image" with the following code:

```
def predict_image(model, img_path):
    model.eval()

    img = Image.open(img_path).convert("RGB")
    # Apply the same transforms as training, and add a batch dimension (unsqueeze)
    img = transform(img).unsqueeze(0).to(device)

    with torch.no_grad():
        # TODO: Get the model prediction
        # 1. Get the raw model outputs (logits)
        output = model(img)

        # 2. Get the predicted class index
        _, pred = torch.max(output, 1)

    return class_names[pred.item()]

# Test the function (this path should exist)
test_img_path = "/content/dataset/paper/0loadHv0B3m7I.png"
prediction = predict_image(model, test_img_path)
print(f"Model prediction for {test_img_path}: {prediction}")
```

The output of the final print statement is shown as:

```
Model prediction for /content/dataset/paper/0loadHv0B3m7I.png: paper
```

boilerplate.ipynb - Colab

https://colab.research.google.com/drive/1CtjhXzFLYRNwAxISlrkvpjPeeg5fLV#scrollTo=wexPK8V3y3fx

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Step 8: Play the Game!

This code is complete. If your model is trained, you can run this cell to have the model play against itself.

```
(28) ✓ On
import random
import os

def pick_random_image(class_name):
    folder = f"/content/dataset/{class_name}"
    files = os.listdir(folder)
    img = random.choice(files)
    return os.path.join(folder, img)

def rps_winner(move1, move2):
    if move1 == move2:
        return "Draw"
    rules = {
        "rock": "scissors",
        "paper": "rock",
        "scissors": "paper"
    }

    if rules[move1] == move2:
        return "Player 1 wins! {move1} beats {move2}"
    else:
        return "Player 2 wins! {move2} beats {move1}"

# -----
# 1. Choose any two random classes
# -----
choices = ["rock", "paper", "scissors"]
c1 = random.choice(choices)
c2 = random.choice(choices)

import math = npick_random_image(rps)
```

Variables Terminal ✓ 5:55PM T4 (Python 3)

boilerplate.ipynb - Colab

https://colab.research.google.com/drive/1CtjhXzFLYRNwAxISlrkvpjPeeg5fLV#scrollTo=wexPK8V3y3fx

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

On
1. Choose any two random classes

choices = ["rock", "paper", "scissors"]
c1 = random.choice(choices)
c2 = random.choice(choices)

img1_path = pick_random_image(c1)
img2_path = pick_random_image(c2)

print("\nRandomly selected images:")
print("Image 1:", img1_path)
print("Image 2:", img2_path)

2. Predict their labels using the model

p1 = predict_image(model, img1_path)
p2 = predict_image(model, img2_path)

print("\nPlayer 1 shows:", p1)
print("Player 2 shows:", p2)

3. Decide the winner

print("\nRESULT:", rps_winner(p1, p2))

*** Randomly selected images:
Image 1: /content/dataset/rock/Au9hY8uo2ixalMKQI.png
Image 2: /content/dataset/scissors/R0gc9adml3hXrDxq.png

Player 1 shows: rock
Player 2 shows: scissors

RESULT: Player 1 wins! rock beats scissors

Variables Terminal ✓ 5:55PM T4 (Python 3)